


# Threads



# What is a Thread?

A thread is a lightweight execution unit inside a process

Multiple Threads share:

- Address space
- File descriptors
- Heap
- Global

Threads have their own:

- Stack
- Registers
- thread ID

# When to Use Threads

- Concurrent I/O (e.g., sensors, logging)
- Background tasks (telemetry, watchdog pings)
- Producer/consumer pipelines (ring buffers)
- Real-time task separation (with priorities & policies)

# Process – Thread Memory

Each process has a separate memory allocated.



A process can have multiple threads inside it.



# Core pthread APIs

Syntax	Parameters	Return & Notes
<pre>int pthread_create(pthread_t *thread,                   const pthread_attr_t *attr,                   void *(*start_routine)(void *),                   void *arg)</pre>	thread: out param; attr: NULL for defaults; start_routine: function; arg: user data	0 success; non-zero error (e.g., EAGAIN). New thread starts at start_routine.
<pre>int pthread_join(pthread_t thread, void **retval)</pre>	thread: target; retval: receives return from pthread_exit() or function return	Blocks until thread terminates. ESRCH/EINVAL/EDEADLK possible.
<pre>void pthread_exit(void *retval)</pre>	retval: return pointer for joiner	Does not return. Cleans up thread resources properly.
<pre>int pthread_detach(pthread_t thread)</pre>	Make thread unjoinable; resources auto-released at exit	Use for fire-and-forget workers. Don't join() after detach.

# Thread ID

- Each thread within a process is uniquely identify by a thread ID.
- This thread ID is returned by the caller of pthread\_create()
- A thread can obtain its thread id by using pthread\_self()

```
#include <pthread.h>  
pthread_t pthread_self(void);
```

# Creating A simple Threads

```
#include <stdio.h>
#include <pthread.h>
#include <unistd.h>

void* incrementFunction(){
for (int i = 0; i < 5; i++) {
    printf("Thread: i = %d\n", i); sleep(1);
}
return NULL; // Thread ends
}

void* decrementFunction(){
for (int j = 5; j > 0; j--) {
    printf("Thread: j = %d\n", j); sleep(1);
}
return NULL; // Thread ends
}
```

```
int main(void){
pthread_t thread_id1, thread_id2;

pthread_create(&thread_id1,NULL,
incrementFunction,NULL);
pthread_create(&thread_id2,NULL,
decrementFunction,NULL);

pthread_join(thread_id1,NULL);
pthread_join(thread_id2,NULL);

}
```

```

#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

void *thread_func_T2(void *arg) {
    printf("T2: Starting work...\n");
    sleep(2);
    printf("T2: Work completed.\n");
    return (void *)"T2 done";
}

void *thread_func_T1(void *arg) {
    pthread_t *t2 = (pthread_t *)arg;
    void *retval;
    printf("T1: Waiting for T2 to finish...\n");
    pthread_join(*t2, &retval);
    printf("T1: T2 finished with message: %s\n", (char *)retval);
    printf("T1: Doing its own work...\n");
    sleep(1);
    printf("T1: Work completed.\n");
    return (void *)"T1 done";
}

void *thread_func_generic(void *arg) {
    int id = *((int *)arg);
    printf("T%d: Doing work...\n", id);
    sleep(1);
    printf("T%d: Work completed.\n", id);
    return (void *)"Generic done";
}

```

```

int main() {
    pthread_t t1, t2, t3, t4;
    int id3 = 3, id4 = 4;
    void *retval;

    // Create T2 first
    if (pthread_create(&t2, NULL, thread_func_T2, NULL) != 0)
        {perror("Failed to create T2"); exit(EXIT_FAILURE);}

    // Create T1 and pass T2's ID to it
    if (pthread_create(&t1, NULL, thread_func_T1, &t2) != 0)
        {perror("Failed to create T1"); exit(EXIT_FAILURE);}

    // Create T3 and T4
    pthread_create(&t3, NULL, thread_func_generic, &id3);
    pthread_create(&t4, NULL, thread_func_generic, &id4);

    // Main thread joins T1, T3, and T4
    pthread_join(t1, &retval);
    printf("Main: T1 finished with message: %s\n", (char *)retval);

    pthread_join(t3, &retval);
    pthread_join(t4, &retval);


    printf("Main: All threads completed. Exiting.\n");
    return 0;
}

```

Aspect	Process	Thread
<b>Basic Definition</b>	A process is an independent program in execution with its own memory space.	A thread is a lightweight unit of execution within a process.
<b>Memory Space</b>	Has its own separate virtual memory space.	Shares the process's memory space (text, data, heap).
<b>Stack</b>	Each process has its own independent stack.	Each thread has its own stack within the same process.
<b>Creation Overhead</b>	High – memory allocation, page tables, resources must be created.	Low – shares most resources, only stack + thread control block created.
<b>Context Switch Cost</b>	Higher – switching memory context and kernel structures.	Lower – same memory space, lighter task switch.
<b>Communication</b>	Requires IPC mechanisms (pipes, shared memory, message queues).	Very fast – communicates via shared memory within the process.
<b>Isolation</b>	Strong isolation – one process crash does not affect others.	Weak isolation – thread crash can terminate the entire process.
<b>Resource Sharing</b>	Not shared by default; must use explicit IPC.	Threads share code, data, heap, file descriptors.
<b>Scheduling Unit in Linux</b>	The process's main thread is scheduled as a task.	Each thread is scheduled as an individual task.
<b>Concurrency</b>	Limited by heavier switch + isolation boundaries.	High concurrency with lightweight switching.
<b>Security</b>	More secure due to isolation.	Less secure because shared address space.
<b>Suitable For</b>	Large independent applications, services, fault isolation.	Parallel tasks within the same app, real-time workflows, worker pools.
<b>Linux Representation</b>	Represented by a task_struct with unique PID.	Represented by task_struct with unique TID but same thread group ID.
<b>Creation APIs</b>	fork(), exec()	pthread_create() or clone() with flags.
<b>Failure Impact</b>	Process crash usually does not affect other processes.	Thread crash can bring down the entire process.
<b>Typical Use in Embedded Linux</b>	Running separate services, daemons, system processes.	Real-time tasks, communication handlers, worker threads, parallel computations.



# Mutex



# Why Synchronization Is Required

When multiple threads run inside the same process, they share memory:

- Global variables
- Static variables
- Heap memory
- File descriptors
- Peripheral buffers (UART, CAN, SPI)

Because of this, two or more threads may try to access or modify the same shared data simultaneously.

# Typical problems that occur without synchronization

## Race Conditions

Two threads modifying the same variable at the same time → unpredictable output.

## Data Corruption

Concurrent writes to shared buffers (e.g., ring buffers, sensor logs) can corrupt data.

## Inconsistent State

Thread A updates part of a structure while Thread B reads it → B gets half-updated data.

## Deadlocks & Priority Inversion

Critical in embedded and real-time systems.


To prevent these problems, we need a mechanism that ensures only one thread at a time can access a shared resource → MUTEX

# Mutex

A mutex (**MUTual EXclusion object**) is a synchronization primitive that allows only one thread to own or lock a resource at a time.

It acts like:

- A key that only one thread can hold.
- Others must wait until the thread releases it.



# How a Mutex Works

- State of a Mutex



## Basic Operations

### 1. `pthread_mutex_lock()`

1. If unlocked → lock it and continue.
2. If locked → thread goes into **blocked state**.

### 2. `pthread_mutex_unlock()`

1. Releases mutex so another waiting thread can acquire it.

### 3. `pthread_mutex_trylock()`

1. Attempts lock without blocking.

# When to Use a Mutex

Shared data must be modified

Updating a global counter  
Writing to shared sensor buffers  
Modifying linked lists, queues, ring buffers  
Accessing shared hardware (I2C/SPI/UART write sequences)

Only one thread must access a critical section at a time

- Logging to a file
- Driving a motor control output
- Updating PWM values
- Shared communication stack (TCP/IP, CAN, UART Rx/Tx queues)

Protecting non-atomic operations

Even simple operations like `counter++` need mutex protection if accessed by multiple threads.

# Sample Example - Mutex

```
#include <pthread.h>
#include <stdio.h>

int shared_counter = 0;
pthread_mutex_t lock;

void* worker(void *arg)
{
    for (int i = 0; i < 100000; i++) {
        pthread_mutex_lock(&lock);
        shared_counter++; // safe update
        pthread_mutex_unlock(&lock);
    }
    return NULL;
}

int main()
{
    pthread_t t1, t2;
    pthread_mutex_init(&lock, NULL);

    pthread_create(&t1, NULL, worker, NULL);
    pthread_create(&t2, NULL, worker, NULL);

    pthread_join(t1, NULL);
    pthread_join(t2, NULL);

    printf("Final counter value: %d\n",
    shared_counter);
    pthread_mutex_destroy(&lock);
    return 0;
}
```

# Best Practices

Minimize time spent inside mutex

Keep critical sections small and fast.

Avoid taking multiple mutexes in different orders

Prevent deadlocks.

Use priority-inheritance mutexes for real-time tasks

Avoids priority inversion issues.

Do not lock in high interrupt frequency handlers

Better use lock-free or atomic operations.

Always unlock before returning or on error paths

Use cleanup handlers if needed.

# Scenario: Robot Battery Monitor

- Imagine you are developing basic software for a home-cleaning robot. The robot continuously moves around the house and its battery slowly drains. When the battery becomes too low, the robot must automatically go to the charging dock and start charging.
- To simulate this behavior, you must write a C program that uses:
  - **One background thread** → simulates the battery draining
  - **Main thread** → checks battery level periodically
  - **fork() + exec()** → to run an external “charging program” (such as your own charger program that prints *“Robot is now charging...”*)
- **Problem Statement:** Design and implement a C program that simulates a robot’s battery monitoring system. The program must meet the following requirements:

# Part 1: Battery Drain Simulation (Thread)

Create a thread named **battery\_thread** which:

- Starts with battery\_level = 100%.
- Every 1 second:
  - Reduces the battery by a random value between **1% and 10%**.
  - Prints: [Sensor] Battery = XX%
- When battery level drops **below 20%**, it must set a shared flag: low\_battery = 1.
- Use a mutex to safely update the shared variables.

# Part 2: Main Thread Monitoring

In the main thread:

- Periodically (every 1 second) check the **low\_battery** flag.
- When **low\_battery == 1**, print: [Main] Low battery detected! Going to charging station...

Then trigger the charging process.

# Part 3: Charging Action

When low battery is detected:

- Use fork() to create a child process.
- In the **child process**, run an external program such as:
  - **Your own program** ./charger  
which prints: Robot is now charging...
- In the **parent process**:
  - Wait for the child to finish using wait().
  - Print: [Main] Charging completed!

# Part 4: Reset & Continue

After charging finishes:

- Reset battery\_level back to **100%**.
- Reset the flag low\_battery = 0.
- Continue monitoring forever.